Containerized or Serverless Learning

Comparing the Prediction Latency of Machine Learning Model Deployments on Containerized and Serverless Environments

Diya Ramasamy | Iris Ma | Krithika Balasubramanyam
Raj Shreyas Penukonda | Yuyu Lai

System Purpose and Functionality

Deployed three different machine learning models on containerized and serverless environments as APIs

Logistic Regression (Scikit-Learn)
➔ Input: Tabular Data
➔ Output: Iris Species

Computer Vision Model (TensorFlow)
➔ Input: Grayscale Image
➔ Output: Clothing Type

NLP Model (TensorFlow)
➔ Input: Wine Review
➔ Output: Type of Wine

Iris Dataset

sepal length	sepal width	petal length	petal width	species
5.1	3.5	1.4	0.2	0
4.9	3.0	1.4	0.2	0
⋮	⋮	⋮	⋮	⋮

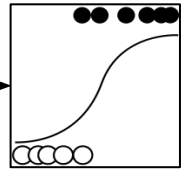

Iris Species

Fashion MNIST

Clothing Type

Wine Review

Aromas include tropical fruit, broom, brimstone and dried herb…

Wine Type

Design

Pipeline

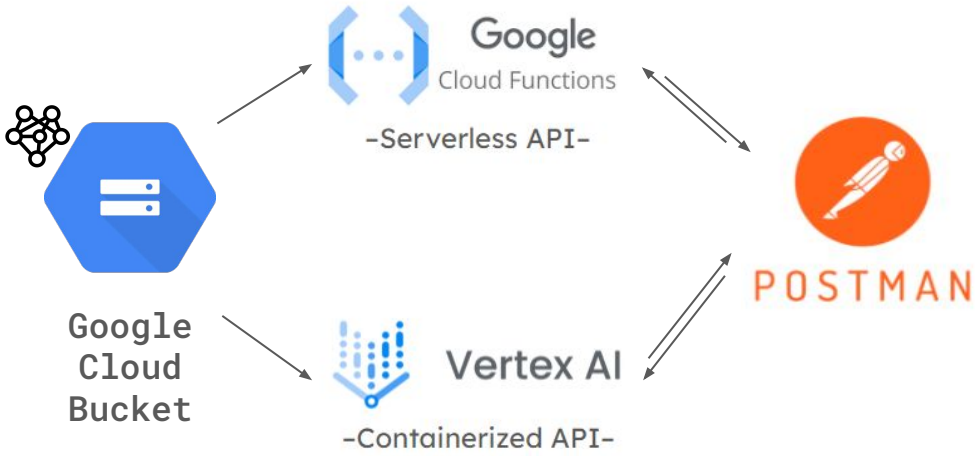

Metrics

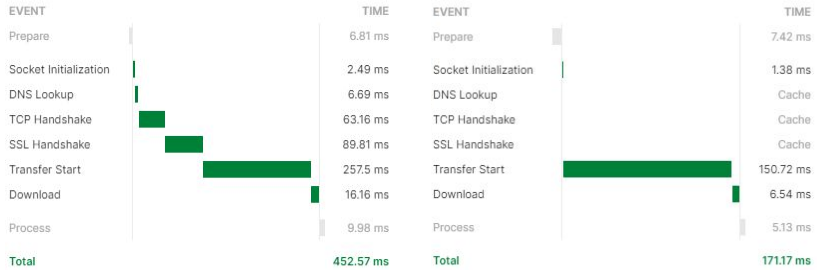

Batch Sizes	Transfer Time (#1)	Transfer Time (#2)	Transfer Time (#3)
1 2 5 10 50 100	…	…	…

Results and Challenges

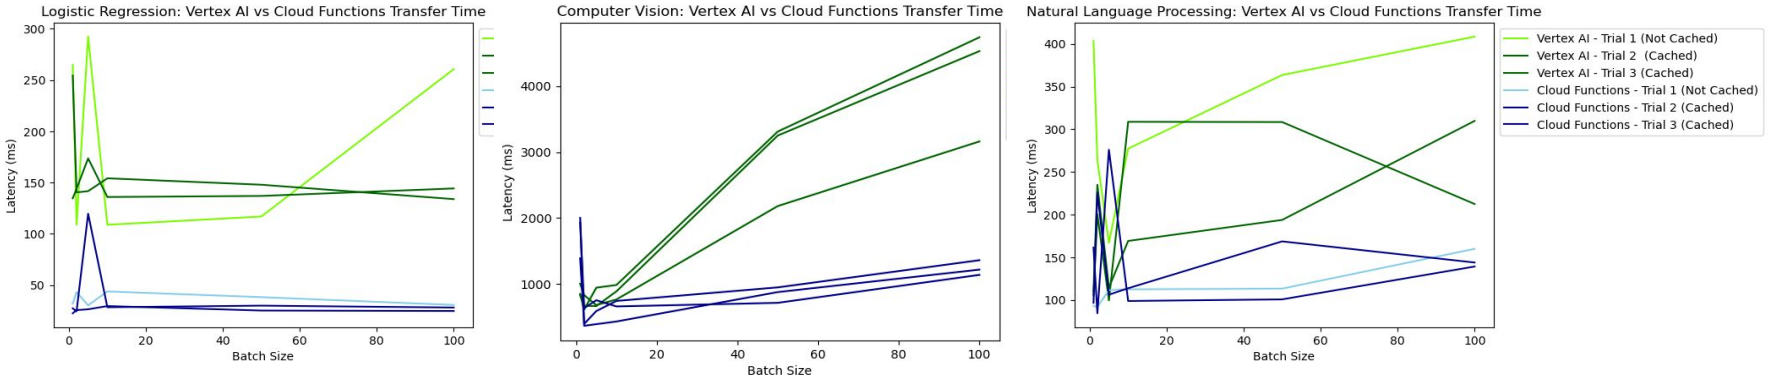

GCF outperforms Vertex AI; dynamic resource allocation allows for consistently low latency

Challenge	**Resolution**
Out of memory issues for large models	Trained another model from scratch with Tensorflow
Framework version conflicts	Trained models with the exact version available on GCP
Authentication issues when calling endpoints	Regenerated tokens every hour